

THEORY ASSIGNMENT 3

Name: Aditya Shivram Mahajan

Roll no: A-10

Reg no: 2017BCS001

Design a MapReduce Application to process the following data for extracting meaningful information from it. Computing Average rating of movies with movie title (Used movies.csv and ratings.csv)

STEP 1: In the first step we will Create a Mapper Class code. The things will be implemented as follows.

MovieNameRatingMapper.java :-

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class MovieNameRatingMapper extends Mapper<LongWritable, Text, IntWritable,
Text> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException,
    InterruptedException {
        // Skip the header row of csv
        if (key.get() == 0 && value.toString().contains("userId")) {
            return;
        }
        String line = value.toString();
        String[] ratingPieces = line.split(",");
        if (ratingPieces.length >= 2) {
            // ratingPieces is: [userId, movieId, rating, timestamp]
            int movieId = Integer.parseInt(ratingPieces[1]);
            String rating = ratingPieces[2];
            // Convert to Hadoop types
            IntWritable mapKey = new IntWritable(movieId);
            Text mapValue = new Text("rating\t" + rating); // Output intermediate
            key,value pair
            context.write(mapKey, mapValue);
        }
    }
}
```

MovieRatingMapper.java :-

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class MovieRatingMapper extends Mapper<LongWritable, Text, IntWritable,
DoubleWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException,
    InterruptedException {
        // Skip the header row of csv
        if (key.get() == 0 && value.toString().contains("userId")) {
            return;
        }
        String line = value.toString();
        String[] ratingPieces = line.split(",");
        if (ratingPieces.length >= 2) {
            // ratingPieces is: [userId, movieId, rating, timestamp]
            int movieId = Integer.parseInt(ratingPieces[1]);
            double rating = Double.parseDouble(ratingPieces[2]); // Convert to
            Hadoop types
            IntWritable mapKey = new IntWritable(movieId);
            DoubleWritable mapValue = new DoubleWritable(rating);
            // Output intermediate key,value pair
            context.write(mapKey, mapValue);
        }
    }
}
```

MovieNameMapper.java :-

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;
public class MovieNameMapper extends Mapper<LongWritable, Text, IntWritable, Text> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException,
```

```

        InterruptedException {
            // Skip the header row of csv
            if (key.get() == 0 && value.toString().contains("movieId")) {
                return;
            }
            String line = value.toString();
            String[] moviePieces = line.split(",");
            if (moviePieces.length >= 2) {
                // ratingPieces is: [movieId,title,genres], want (movieId, title)
                int movieId = Integer.parseInt(moviePieces[0]);
                String title = moviePieces[1]; // Convert to Hadoop types (mark this as
                // the "title" value)
                IntWritable mapKey = new IntWritable(movieId);
                Text mapValue = new Text("title\t" + title);
                // Output intermediate key,value pair
                context.write(mapKey, mapValue);
            }
        }
    }
}

```

STEP 2: Now in the second step we will Create a Reducer Class code. The implementation is as following.

MovieNameRatingReducer.java :-

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class MovieNameRatingReducer extends Reducer<IntWritable, Text, IntWritable, Text>
{
    @Override
    public void reduce(IntWritable key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException{
        double avgRating = 0;
        int numValues = 0;
        String title = "";
        // Add up all the ratings
        for (Text value: values) {
            // Separate the value from the tag
            String parts[] = value.toString().split("\t"); if (parts[0].equals("title")) {
                // Get the title
                title = parts[1];
            } else {
                // Get a rating for this title
            }
        }
    }
}

```

```

        double ratingValue = Double.parseDouble(parts[1]);
        avgRating += ratingValue;
        numValues++;
    }
}
// Divide by the number of ratings to get the average for this movie
avgRating /= numValues;
String outputValue = title + "\n" + avgRating;
Text output = new Text(outputValue);
// Output movieId, title and avg rating
context.write(key, output);
}
}

```

MovieRatingReducer.java :-

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.mapreduce.Reducer;public class MovieRatingReducer extends
Reducer<IntWritable, DoubleWritable, IntWritable,
DoubleWritable> {
    @Override
    public void reduce(IntWritable key, Iterable<DoubleWritable> values, Context context)
    throws IOException, InterruptedException{
        double avgRating = 0;
        int numValues = 0;
        // Add up all the ratings
        for (DoubleWritable value: values) {
            avgRating += value.get();
            numValues++;
        }
        // Divide by the number of ratings to get the average for this movie
        avgRating /= numValues;
        DoubleWritable average = new DoubleWritable(avgRating);
        // Output the average rating for this movie
        context.write(key, average);
    }
}
}

```

STEP 3: Now we Create Driver code

MovieNamesRatings.java :-

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.MultipleInputs;
public class MovieNamesRatings {
    public static void main(String[] args) throws Exception {
        // Check that the input and output path have been provided
        if (args.length != 3) {
            System.err.println("Syntax: MovieNamesRatings <ratings input path>
            <titles input path>
            <output path>");
            System.exit(-1);
        }
        // Create an instance of the MapReduce job
        Job job = new Job();
        job.setJarByClass(MovieNamesRatings.class);
        job.setJobName("Average movie rating by movie title");
        // Set input and output locations
        MultipleInputs.addInputPath(job, new Path(args[0]), TextInputFormat.class,
        MovieNameRatingMapper.class);
        MultipleInputs.addInputPath(job, new Path(args[1]), TextInputFormat.class,
        MovieNameMapper.class);
        FileOutputFormat.setOutputPath(job, new Path(args[2]));
        // Set reducer classes
        job.setReducerClass(MovieNameRatingReducer.class);
        // Set output key/value
        job.setOutputKeyClass(IntWritable.class);
        job.setOutputValueClass(Text.class);
        // Run the job and then exit the
        programSystem.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

STEP 4: Compile all JAVA files using -classpath option in javac

```
javac -cp
$HADOOP_INSTALL/share/hadoop/mapreduce/hadoop-mapreduce-client-core-
2.9.2.jar:$HADOOP_INSTALL/share/hadoop/common/hadoop-common-2.9.2.jar -d
<OUTPUT_PATH_FOR_COMPILED_FILES> <INPUT_JAVA_FILES_PATH>/*.java
```

```
=====

skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce$ javac -cp
/usr/skystone/hadoop/hadoop-2.9.2/ share/hadoop/mapreduce/hadoop-mapreduce-
client-core-2.9.2.jar:/usr/local/hadoop/hadoop-2.9.2/share/hadoop/common/hadoop-common-2.9.2.jar -d
/mnt/E4687B61687B3182/map-reduce/src *.java
```

Note: Some input files use or override a deprecated API.

Note: Recompile with -Xlint :deprecation for details.

```
skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce$
```

STEP 5: Create JAR file

```
jar -vcf <PATH_TO_INITIALIZED_NEWLY_CREATED_JAR_FILE>/mr.jar
<INPUT_PATH_OF_ALL_COMPILED_CLASS_FILES>/*.class
```

```
=====

skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce$ jar -vcf mr.jar *.class
added manifest
adding: movieNamemapper.class(in = 2272) (out= 958) (deflated 57%)
adding: movieNameRatingmapper.class(in = 2284) (out= 961) (deflated 57%)
adding: movieNameRatingReducer.class(in = 2472) (out= 1857) (deflated 57%)
adding: movieNamesRatings.class(in = 1638) (out= 878) (deflated 46%)
adding: movieRatingmapper.class(in = 1957) (out= 819) (deflated 58%)
adding: movieRatingReducer.class(in = 1785) (out= 781) (deflated 58%)
adding: movieRatings.class(in = 1414) (out= 885) (deflated 43%)
skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce$
```

STEP 6: Run the JAR file

```
$HADOOP_INSTALL/bin/hadoop jar <PATH_OF_JAR_FILE>/mr.jar
<PATH_OF_DRIVER_CLASS>/MovieNamesRatings
<INPUT_PATH_TO_RATING.CSV_FILE_OR_DATA>/ratings.csv
<INPUT_PATH_TO_MOVIES>CSV_FILE>/movies.csv
<OUTPUT_PATH_WHERE_OUTPUT_IS_CREATED>
```

=====

skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce\$
/usr/local/hadoop/hadoop-2.9.2/bin/hadoop jar mr.jar MovieNamesRatings
/mnt/E4687B61687B3182/map-reduce/data/ratings.csv
/mnt/E4687B61687B3182/map-reduce/data/movies.csv /mnt/E4687B61687B3182/map-reduce/OUTPUTS

06/12/12 12:11:58 INFO mapreduce.Job: counters: 38

File System Counters

FILE: Number of bytes read-7886222
FILE: Number of bytes written-5352124
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0

map-Reduce Framework

map input records-188837
map output records=133836
map output bytes-1218832
map output materialized bytes-1411718
Input split bytes-133
Combine input records=3
Combine output records=3
Reduce input groups-9724
Reduce shuffle bytes-1411718
Reduce input records-188836
Reduce output records-9724
Spilled Records-281672
Shuffled maps
Failed Shuffles=3
merged map outputs-1
GC time elapsed
Total committed heap usage (bytes) =488247888

Shuffle Errors

BAD ID=0
CONNECTION=0
IO ERROR=0
WRONG LENGTH=0
WRONG
WRONG REDUCE=0

File Input Format Counters

Bytes Read-2483723

File Output Format Counters

Bytes Written-137486

STEP 7: Display Output File

```
$HADOOP_INSTALL/bin/hdfs dfs -cat <PATH_OF_OUTPUT_FILE>/part-r-00000
```

```
=====
```

```
skystone@skystone-HP-Notebook:/mnt/E4687B61687B3182/map-reduce$
```

```
/usr/local/hadoop/hadoop-2.9.2/bin/hdfs dfs -cat
```

```
/mnt/E4687B61687B3182/map-reduce/OUTPUT9/part-r-00000
```

```
=====
```

8593 Juice (1992)

4.0

8596

Revenge of the Pink Panther (1978)

3.25

8600

Angels with Dirty Faces (1938)

3.75

8601

Zero de conduite (Zero for Conduct) (Zéro de conduite: Jeunes diables au collège) (1933)

4.0

8604

Taxi (1998)

3.2

8605

Taxi 3 (2003)

3.25

8607

Tokyo Godfathers (2003)

3.9

8609

Our Hospitality (1923)

3.25

8610


All of Me (1984)

2.5

8611 "Farmer's Daughter

3.5

8614



Overboard (1987)
3.0
8617
Butterfield 8 (1960)
3.0
8620 "Exterminating Angel
1.5
8622
Fahrenheit 9/11 (2004)3.4864864864864864
8623
Roxanne (1987)
3.272727272727273
8626
Dr. Terror's House of Horrors (1965)
1.5
8632
Secret Society (2002)
0.5
8633 "Last Starfighter
3.5714285714285716
8636
Spider-Man 2 (2004)
3.8037974683544302
8638
Before Sunset (2004)
3.7
8640
King Arthur (2004)
2.9615384615384617
8641
Anchorman: The Legend of Ron Burgundy (2004)
3.7719298245614037
8643 "Cinderella Story
3.2777777777777777
8644 "I
3.4918032786885247
8645 "Maria Full of Grace (Maria
3.9375
8650
Long Day's Journey Into Night (1962)
3.5
8656 "Short Film About Killing



3.75

8665 "Bourne Supremacy

3.7866666666666666

8666

Catwoman (2004)

1.3333333333333333

8667

A Home at the End of the World (2004)3.5

8670 "Testament of Dr. Mabuse

4.0

8677

Flash Gordon Conquers the Universe (1940)

1.25

8684 "Man Escaped

3.5

8685 "Miracle of Marcelino

3.5

8690

Slaughterhouse-Five (1972)

4.0

8695 "Bachelor and the Bobby-Soxer

3.0