


STEP1:CREATE DATA FILE

FIRST create the data file named as file1.txt



STEP2: Start HDFS (Namenode and Datanode) and YARN (Resource Manager and Node Manager)

Run following commands

```
C:\Users\Mansi>cd c:\hadoop
c:\hadoop>sbin\start-dfs
c:\hadoop>sbin\start-yarn
starting yarn daemons..
```

Namenode, Datanode, Resource Manager and **Node Manager** will be started in few minutes and

3) STEPS 3

- Create a directory (say 'input') in HDFS to keep all the text files (say 'file1.txt') to be used for counting words.

```

Command Prompt
Microsoft Windows [Version 10.0.17134.1006]
(c) 2018 Microsoft Corporation. All rights reserved.

C:\Users\Mansi>cd c:\hadoop
c:\hadoop>bin\hdfs dfs -mkdir input

```

- Copy the text file(say 'file1.txt') from local disk to the newly created 'input' directory in HDFS.

```

Command Prompt
Microsoft Windows [Version 10.0.17134.1006]
(c) 2018 Microsoft Corporation. All rights reserved.

C:\Users\Mansi>cd c:\hadoop
c:\hadoop>bin\hdfs dfs -copyFromLocal c:/file1.txt input

```

- Run the wordcount job provided in
%HADOOP_HOME%\share\hadoop\mapreduce\hadoop-mapreduce-examples-2.2.0.jar
Command :
C:\hadoop>bin\yarn jar share/hadoop/mapreduce/hadoop-mapreduce-examples-2.2.0.jar
wordcount input output

Output of command line after executing above command:

```

13:22:02 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
13:22:03 INFO input.FileInputFormat: Total input paths to process : 1
13:22:03 INFO mapreduce.JobSubmitter: number of splits:1
:
:
13:22:04 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1391412385921_0002
13:22:04 INFO impl.YarnClientImpl: Submitted application application_1391412385921_0002 to
ResourceManager at /0.0.0.0:8032
13:22:04 INFO mapreduce.Job: Running job: job_1391412385921_0002
13:22:14 INFO mapreduce.Job: Job job_1391412385921_0002 running in uber mode : false

```

13:22:14 INFO mapreduce.Job: map 0% reduce 0%
13:22:22 INFO mapreduce.Job: map 100% reduce 0%
13:22:30 INFO mapreduce.Job: map 100% reduce 100%
13:22:30 INFO mapreduce.Job: Job job_1391412385921_0002 completed successfully
13:22:31 INFO mapreduce.Job: Counters: 43

File System Counters

FILE: Number of bytes read=89
FILE: Number of bytes written=160142
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=171
HDFS: Number of bytes written=59
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=5657
Total time spent by all reduces in occupied slots (ms)=6128

Map-Reduce Framework

Map input records=2
Map output records=7
Map output bytes=82
Map output materialized bytes=89
Input split bytes=116
Combine input records=7
Combine output records=6
Reduce input groups=6
Reduce shuffle bytes=89
Reduce input records=6
Reduce output records=6
Spilled Records=12
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=145
CPU time spent (ms)=1418
Physical memory (bytes) snapshot=368246784

Virtual memory (bytes) snapshot=513716224
Total committed heap usage (bytes)=307757056
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=55
File Output Format Counters
Bytes Written=59

MAPREDUCE CODES:

MAPPER:

```
package WordCount.WordCount;

import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WC_Mapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable> output,
        Reporter reporter) throws IOException{
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()){
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

```
}  
  
}
```

REDUCER:

```
import java.io.IOException;  
import java.util.Iterator;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapred.MapReduceBase;  
import org.apache.hadoop.mapred.OutputCollector;  
import org.apache.hadoop.mapred.Reducer;  
import org.apache.hadoop.mapred.Reporter;  
  
public class WC_Reducer extends MapReduceBase implements  
Reducer<Text,IntWritable,Text,IntWritable> {  
    public void reduce(Text key, Iterator<IntWritable> values,OutputCollector<Text,IntWritable>  
output,  
    Reporter reporter) throws IOException {  
        int sum=0;  
        while (values.hasNext()) {  
            sum+=values.next().get();  
        }  
        output.collect(key,new IntWritable(sum));  
    }  
}
```

Wordcount runner program

```
import java.io.IOException;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapred.FileInputFormat;  
import org.apache.hadoop.mapred.FileOutputFormat;  
import org.apache.hadoop.mapred.JobClient;  
import org.apache.hadoop.mapred.JobConf;  
import org.apache.hadoop.mapred.TextInputFormat;  
import org.apache.hadoop.mapred.TextOutputFormat;  
public class WC_Runner {  
    public static void main(String[] args) throws IOException{  
        JobConf conf = new JobConf(WC_Runner.class);
```

```

        conf.setJobName("WordCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
        conf.setReducerClass(WC_Reducer.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf,new Path(args[0]));
        FileOutputFormat.setOutputPath(conf,new Path(args[1]));
        JobClient.runJob(conf);
    }
}

```

CHECK OUTPUT :



```

C:\Users\Mansi>cd c:\hadoop

c:\hadoop>bin\hdfs dfs -cat output/*
movieId 1
title 1
genres 1
Toy 1
Story 1
Adventure 2
Animation 1
Children 2
Comedy 2
Fantasy 2
Jumanji 1
Grumpier 1
Old 1
Men 1
Romance 1

```

Conclusion: I have extracted the Count of each word in the file given named file1.txt